

Double-UV Relative-Momentum Enhancement in the GL00 $\gamma\gamma \rightarrow \gamma\gamma$ Integrand

Investigation note

June 14, 2026

Abstract

This note documents the current understanding of a large-weight feature seen in the generated GL00 two-loop integrand for $\gamma\gamma \rightarrow \gamma\gamma$ with the nested integrated UV counterterm included. The issue is not an ordinary uncanceled UV divergence in a generic double-UV direction. The problematic samples lie close to a relative-momentum ridge: two sampled loop-momentum-basis (LMB) rays become hard, nearly collinear, and have comparable radii, so that one physical internal momentum is a small difference of the two hard basis momenta. The resulting enhancement is locally integrable, but it creates large Monte-Carlo variance because the current sampling channels do not parameterize that relative momentum directly. LMB multichanneling is not merely a random choice of LMB: it includes a normalized OSE-product prefactor, and that prefactor strongly suppresses inappropriate channels near the edge-small limit. The remaining question is how much this discrete reweighting, and in particular its exponent α , can reduce the variance without an additional continuous relative-momentum map. A fixed-ray option was added to the UV profiler to make this non-generic direction visible in profiling runs.

1 Executive summary

For the regenerated all-LMB setup, GAMMALOOP reports all eleven loop momentum bases as active channels:

$$(5, 10), (4, 10), (5, 9), (4, 9), (5, 8), (4, 8), (5, 7), (4, 7), (5, 6), (4, 6), (4, 5).$$

This required one small display/introspection correction: the display command must read the generated multichannel setup instead of recomputing the old heuristic.

The short all-LMB Monte-Carlo run still found its largest signed evaluations in channels (5,10) and (5,9), not in a channel that explicitly uses edge 4 as a basis variable. That is consistent with a variance problem rather than a pure channel-enumeration problem. In the original graph-generation basis,

$$e_4 = K_0, \tag{1}$$

$$e_5 = -K_0 + K_1 + P_0, \tag{2}$$

$$e_9 = K_1, \tag{3}$$

$$e_{10} = -P_2 + K_1 + P_0 + P_1. \tag{4}$$

In channel (5,10), where $u = e_5$ and $v = e_{10}$ are sampled directly,

$$e_4 = v - u + P_2 - P_1. \tag{5}$$

Thus a sample with u and v both hard, nearly collinear, and of similar radius maps to a much smaller edge-4 momentum. This is precisely the geometric configuration that the OSE channel weights can identify only indirectly: they can favor LMBs containing edge 4, but they do not replace the continuous variables (u, v) by common and relative momenta.

2 What the inspected weight contains

The approach plots below use the value printed by GAMMALOOP as “The evaluation of integrand”. This is the event weight passed to the integrator. The same inspect output also prints the parameterization jacobian separately and a much smaller “integrand result” before multiplication by that jacobian. Therefore the plotted quantity already includes the spherical parameterization jacobians and the conformal-map jacobians. The large-radius slope in the plots is a statement about the actual weight seen by the integrator, not merely about the unweighted analytic integrand.

3 Minimal toy model

The minimal structure needed to reproduce the observed mechanism is not a full two-loop numerator. It is enough to keep:

- two hard sampled LMB momenta u and v ;
- a physical internal momentum $q = v - u + Q$ that is a difference of those hard momenta plus a fixed external shift Q ;
- the post-subtraction and parameterization effect that leaves a positive power of the common hard radius in the event weight;
- an energy denominator sensitive to the small relative momentum.

The toy event weight used for the figures is

$$W_{\text{toy}}(u, v) = C(\hat{u}, \hat{v}, \rho) \frac{R^2}{E_q}, \quad E_q = \sqrt{|v - u + Q|^2 + q_0^2}, \quad R = \frac{|u| + |v|}{2}, \quad (6)$$

where $\rho = |v|/|u|$ and C is a smooth bounded angular factor. The constant q_0 regulates the exactly coincident toy edge and stands in for the fact that the full GL00 expression has masses, external shifts, and numerator cancellations. The power R^2 should not be read as a universal UV degree. It is the minimal effective power required to match the observed weighted large-radius approach after all jacobians are included.

Let the relative separation scale as

$$|v - u + Q| \simeq \epsilon R \quad (7)$$

for a fixed but small opening/radius mismatch ϵ . Then

$$W_{\text{toy}} \simeq \frac{R^2}{\epsilon R} = \frac{R}{\epsilon}. \quad (8)$$

This gives a slope-one growth along a ray where the two UV momenta are scaled together while their relative mismatch is held fixed. If the two hard momenta coalesce more quickly than $1/R$, the toy model crosses over toward R^2/q_0 . For generic separated directions the coefficient is much smaller; the bad region is the ridge where $\rho \simeq 1$ and $\theta(u, v) \simeq 0$.

The singularity is locally integrable in the relative momentum. In relative variables $d^3q = |q|^2 d|q| d\Omega_q$, a $1/|q|$ factor gives

$$d^3q \frac{1}{|q|} \sim |q| d|q| d\Omega_q,$$

so the local integral is finite. Monte Carlo variance can nevertheless be large, especially when the sampling variables are the hard momenta u and v rather than the common momentum and the relative momentum q .




Figure 1: Toy-model radial scaling. For fixed nonzero relative mismatch ϵ , the weighted integrand approaches a slope-one growth, $W_{\text{toy}} \sim R/\epsilon$. Exact coalescence crosses over to slope two in this deliberately simplified model.




Figure 2: Toy relative-momentum ridge at fixed large common radius. The enhancement is concentrated near equal radius, $\rho = 1$, and small opening angle. Existing LMB channels can reduce coefficients, but they do not create a direct sampling variable for this ridge.

4 Standalone proxy integrand experiment

I also built a standalone Python mock-up that is closer to the actual subtracted-integrand mechanism than the purely phenomenological toy model above: `DoubleUVIntegrableSingularity/mock_gl00_integrand_proxy.py`. It defines an original relative-edge factor and the corre-

sponding local double-UV counterterm obtained from the 3D expansion that drops the fixed external shift in the relative edge:

$$W_{\text{orig}}^{(K)}(u, v) = N_K(u, v) \frac{[(E_u + E_v)/2]^3}{E_u E_v} \frac{1}{E(v - u + Q, m_{\text{phys}})}, \quad (9)$$

$$W_{\text{UV}}^{(K)}(u, v) = N_K(u, v) \frac{[(E_u + E_v)/2]^3}{E_u E_v} \frac{1}{E(v - u, m_{\text{UV}})}, \quad (10)$$

$$W_{\text{sub}}^{(K)}(u, v) = W_{\text{orig}}^{(K)}(u, v) - W_{\text{UV}}^{(K)}(u, v). \quad (11)$$

Here u and v are the two hard sampled LMB momenta and Q is the fixed external shift in the physical relative edge. The optional integrated-CT switch only changes the smooth numerator coefficient N_K : with the switch off it contains the original numerator proxy, while with the switch on it adds a smooth integrated-self-energy proxy. It does not introduce a new denominator.

The same script also contains a second mode that mimics doing the local UV Taylor expansion in four dimensions first and only then taking the CFF expression of the resulting massive vacuum topology. The relative factor in that mode is

$$F_{\text{vac}}(q; R, m) = \frac{1}{E(q, m) [2R + E(q, m)]},$$

with the UV counterterm using a finite vacuum mass $m_{\text{UV}} = 9.188$. This is not intended to be a full CFF formula, but it keeps the two features relevant to the present question: the relative propagator is a vacuum object before CFF, and after CFF it is represented by a massive on-shell energy.

The fitted large-radius slopes of $|W_{\text{sub}}^{(K)}|$ are:

mode	ray choice	without integrated CT	with integrated CT
3D expansion	locked equal-norm collinear rays	+1.000	+1.000
3D expansion	parallel rays with 4% radius mismatch	-1.000	-1.000
3D expansion	equal radius with 1° opening angle	-1.000	-1.000
3D expansion	generic directions	-1.000	-1.000
4D first, then CFF	locked equal-norm collinear rays	+1.000	+1.000
4D first, then CFF	parallel rays with 4% radius mismatch	-1.002	-1.002
4D first, then CFF	equal radius with 1° opening angle	-0.997	-0.997
4D first, then CFF	generic directions	-1.000	-1.000

This reproduces the qualitative GL00 behavior: the subtracted proxy is UV safe for generic double-UV rays and for fixed nonzero angular/radial mismatch, but it grows when the two hard rays are locked so that $v - u$ remains small while the common radius grows. The same behavior is present with the integrated-CT proxy disabled, so the mechanism is not created by the integrated counterterm; the integrated counterterm only changes the smooth coefficient multiplying the same relative-edge structure. Importantly, the locked-ray growth is still present in the 4D-first proxy. This means the non-uniform double-UV region is not purely an artifact of differentiating the already-3D on-shell energy $E(v - u + Q)$.




Figure 3: Standalone proxy radial scan. In both the direct 3D-expansion mode and the 4D-first vacuum-CFF mode, only exactly locked equal-norm rays grow with the common UV radius. A fixed opening angle or a fixed radius mismatch restores the expected decreasing asymptotic behavior.

At fixed hard common radius, the direct 3D proxy gives

$$|W_{\text{sub}}^{(K)}| \sim \frac{1}{|v - u|}$$

over the relative-momentum range above the small UV regulator. Multiplication by the three-dimensional relative measure gives $|v - u|^2 |W_{\text{sub}}^{(K)}| \sim |v - u|$, so the ridge is locally integrable even though it is large enough to create Monte-Carlo variance. The 4D-first proxy differs here: because the vacuum relative propagator carries a finite UV mass before CFF, the fixed-radius relative scan is flat as $v - u \rightarrow 0$ rather than $1/|v - u|$. Thus doing the derivation in 4D first can remove the literal small-relative-momentum pole introduced by a direct 3D on-shell-energy expansion. It does not, however, remove the radial non-uniformity of the double-UV Taylor expansion on the subspace where the relative four-momentum is not hard.




Figure 4: Standalone proxy relative-momentum scan at fixed common hard radius. The direct 3D expansion has a $1/|v - u|$ ridge. The 4D-first vacuum-CFF proxy is regularized by the massive vacuum propagator, but the three-dimensional relative measure still makes both cases locally integrable.

4.1 Removing the ridge by parameterization

The same proxy script now performs an explicit Monte-Carlo integration over a finite local domain,

$$I = \int d \log R \int_{|q| < q_{\max}} d^3 q W_{\text{sub}}^{(K)}(P = R\hat{n}, q), \quad q = v - u,$$

with $q_{\max} = 0.2$ and 2.5×10^5 samples per channel. The naive channel samples q uniformly in the three-dimensional ball. The singularity-adapted relative channel samples the radial variable with

$$p_r(r) = \frac{2r}{q_{\max}^2}, \quad p_{\text{ridge}}(q) = \frac{1}{2\pi q_{\max}^2 |q|}.$$

Because the direct 3D proxy has $W_{\text{sub}} \sim 1/|q|$, the estimator $W_{\text{sub}}/p_{\text{ridge}}$ is finite as $q \rightarrow 0$. I also tested a conservative balance-heuristic multichannel density,

$$p_{\text{multi}}(q) = 0.25 p_{\text{uniform}}(q) + 0.75 p_{\text{ridge}}(q),$$

which keeps a regular uniform component for the non-singular part of the domain.

This finite-domain experiment should be interpreted carefully. A $1/|q|$ ridge in three relative dimensions is locally integrable, and its local second moment under a flat finite ball is also finite because $\int q^2 dq |1/q|^2$ is finite at $q = 0$. The real integration problem is not just that local endpoint. It is the combination of this thin relative ridge with the hard common radius, the conformal-map Jacobian, and finite statistics. In the unbounded problem the second moment can still be effectively or genuinely divergent if the radial tail is not sampled with a compatible

density. The variance reductions below therefore demonstrate that the relative-momentum channel removes the local rare-event ridge; they are not a proof that the full UV-tail variance is finite.

integrated CT	channel	estimate	relative error	variance reduction	max $ w $
off	uniform	-2.4975×10^5	0.355%	1.0	2.07×10^8
off	ridge	-2.4836×10^5	0.012%	883.9	2.89×10^5
off	multichannel	-2.4840×10^5	0.030%	137.7	3.84×10^5
on	uniform	-3.3226×10^5	0.136%	1.0	2.39×10^7
on	ridge	-3.3177×10^5	0.012%	129.8	3.86×10^5
on	multichannel	-3.3175×10^5	0.030%	20.3	5.15×10^5

This is the desired behavior: the relative-momentum channel removes the integrable singularity from the estimator and collapses the rare-event tail. For example, without the integrated-CT proxy, the maximum absolute sampled weight drops from 2.07×10^8 to 2.89×10^5 , and the variance drops by a factor 884. The conservative multichannel mixture still reduces the variance by a factor 138. The same conclusion holds with the integrated-CT proxy enabled, although the smooth numerator changes the relative coefficient and therefore the exact reduction factor.




Figure 5: Monte-Carlo variance reduction in the standalone proxy. The fixed-hard-radius panel isolates the relative-momentum singularity and shows the large gain from a ridge-adapted channel. The log-hard-radius panel keeps the common-radius growth; the relative channel still helps, but the remaining variance is then dominated by the radial UV tail and would require a complementary hard-radius importance map.

When the same experiment is run in the 4D-first vacuum-CFF proxy, the ridge-adapted channel does not help at fixed hard radius: the relative scan is already flat as $q \rightarrow 0$, so oversampling small q only wastes samples. This is a useful cross-check. The relative channel should be attached to the direct 3D-CFF local counterterm structure where the $1/|k-l|$ ridge is actually present, and combined with ordinary channels through a balance heuristic.

5 GL00 approach data

The GL00 radial scan starts from the high-weight point

$$x_A = (0.9938195811740000, 0.6303276140313304, 0.6037887485521857, \\ 0.9939524818481166, 0.6277285639451646, 0.6151209153931045).$$

At $\lambda = 1$, the two radial variables correspond to approximately

$$r_0 \simeq 4.8240 \times 10^4, \quad r_1 \simeq 4.9307 \times 10^4.$$

The two radii are therefore both in the UV and within about two percent of each other. The angular coordinates are also close; the fitted opening angle in the angular scan is

$$\theta_A \simeq 2.82 \times 10^{-2}.$$

The large-radius fit of the GAMMALOOP inspected weight gives:

case	fit range	slope in $\log W $ vs. $\log \lambda$	points
LMB multichannel off	$\lambda \geq 1$	1.00005	27
LMB multichannel on	$\lambda \geq 1$	1.00424	27

The coefficient changes strongly when LMB multichanneling is enabled, but the large-radius power does not. This is the important observation: the current multichanneling can reduce the normalization of the problematic ray, but it does not alter the radial power seen by the integrator.




Figure 6: GAMMALOOP inspected event weight along the high-weight double-UV ray. All parameterization jacobians are included. The large-radius behavior is approximately linear in the common UV scale.




Figure 7: Compact replot of the same radial data, with large-radius fit slopes. The fit is not intended to model the low-radius region.




Figure 8: Angular approach scan around the same GL00 point. The coefficient is highly sensitive to the relative direction of the two hard rays, supporting the interpretation as a relative-momentum ridge rather than a generic double-UV failure.

6 Why the standard UV profile misses the enhancement

The ordinary `profile ultra-violet` command samples one random three-momentum per LMB loop variable and then rescales the selected subset. It therefore measures the generic large-momentum degree of divergence for that sampled ray. The relative-momentum ridge discussed above is a correlated configuration: two UV rays must be nearly parallel and have nearly equal radii, so that a physical edge momentum such as $v - u + Q$ stays anomalously small compared to the two hard sampled basis momenta. A random UV profile ray hits this condition with probability zero.

There is also a conceptual difference between the UV profiler and the integration-weight plots. The profiler works directly in momentum space and multiplies by the UV scaling measure factor used for the degree-of-divergence test. It does not include the full x -space parameterization jacobian, adaptive discrete-channel probability, or the conformal-map density that define the event weight seen by the integrator. This is why the standard UV profile can be clean while the integrator still sees a high-variance, locally integrable ridge.

To make the correlated direction testable, the profiler now accepts fixed UV ray data:

```
profile ultra-violet ... \
  --uv-ray-directions dx0 dy0 dz0 [dx1 dy1 dz1 ...] \
  --uv-ray-norms n0 [n1 ...]
```

One direction or norm is repeated for all loop variables; otherwise one three-vector and one norm can be supplied per loop variable. If directions are given but no norm is supplied, the default starting norm is E_{cm} .

The same profiler was then run on the fixed ray matching the angular direction and radius ratio of the GL00 approach scan:

```
--uv-ray-directions
-0.66816740923660423 -0.71446756114601095 0.20757749710437134
-0.67621171790860379 -0.69980455265688968 0.23024183078620908
--uv-ray-norms 1 1.0221114989619468
```

The summed, not-per-orientation UV profiles are:

profile	entries	worst slope	r^2	bad entries	arb retry
random ray, all subsets	33	-0.997272	0.999997	0/33	0/33
tailored ray, double-free subsets	11	-0.812993	0.986953	11/11	11/11
tailored ray, single-free subsets	22	-1.000429	1.000000	0/22	0/22

Here “bad” means a fitted slope above the profiler warning threshold -0.9 . The random-ray result is the expected generic UV profile. With the tailored ray, all double-free subsets move to the same fitted behavior,

$$\log |I_{\text{prof}}| \simeq -0.812993 \log s + \text{const}, \quad r^2 = 0.986953,$$

and all of those double-free fits require the arbitrary-precision retry. Single-free subsets remain ordinary.

A follow-up scan continuously shrinks the plot-ray mismatch. Let $a = 1$ denote the GL00 plot ray above and $a = 0$ denote the exactly collinear, equal-radius ray with the first plot-ray direction used for both loop variables. The second direction is

$$\hat{v}(a) = \frac{(1-a)\hat{u} + a\hat{v}_{\text{plot}}}{|(1-a)\hat{u} + a\hat{v}_{\text{plot}}|}, \quad \rho(a) = 1 + a(\rho_{\text{plot}} - 1).$$

The table reports the summed double-free UV-profile slope. All 11 double-free LMB subsets give the same number for each row.

a	$\theta(\hat{u}, \hat{v})$	$\rho - 1$	$10^3 \rightarrow 10^5$	$10^3 \rightarrow 10^7$
1.000	2.817×10^{-2}	2.211×10^{-2}	-0.814	-0.931
0.500	1.408×10^{-2}	1.106×10^{-2}	-0.684	-0.881
0.200	5.633×10^{-3}	4.422×10^{-3}	-0.434	-0.776
0.100	2.817×10^{-3}	2.211×10^{-3}	-0.182	-0.661
0.050	1.408×10^{-3}	1.106×10^{-3}	+0.104	-0.516
0.020	5.633×10^{-4}	4.422×10^{-4}	+0.471	-0.289
0.010	2.816×10^{-4}	2.211×10^{-4}	+0.699	-0.101
0.005	1.408×10^{-4}	1.106×10^{-4}	+0.860	+0.092
0.000	0	0	+1.000	+1.000

A sparse extension to 10^9 gives slopes -0.964 for $a = 1$, -0.820 for $a = 0.1$, -0.611 for $a = 0.02$, -0.376 for $a = 0.005$, and $+1.000$ for $a = 0$. This is the expected non-uniform limit. For any fixed nonzero angular/radius mismatch, probing deep enough eventually resolves the relative momentum and pushes the fit back toward the generic negative UV behavior. As the mismatch is reduced, the crossover moves to larger scales, and over the finite UV windows relevant to the integrator the fitted slope degrades through zero and becomes positive. The exactly coincident GL00-direction ray remains positive because the relative hard momentum is kept fixed to the external shift instead of growing with the common UV scale.

The same two profiler runs were also repeated with `--per-orientation`. For the random ray, the double-free per-orientation profile remained clean: the worst double-free orientation slope was -0.988527 , with no arbitrary-precision retries and no slope above -0.9 . For the tailored ray, the issue is not isolated to a unique orientation. Among the 62 fitted orientations, 48 have at least one double-free entry above -0.9 . The strongest orientations are:

orientation	max slope	avg. double-free slope	bad entries	arb retry
0000+++--+	-0.614883	-0.614883	11/11	11/11
0000--++--	-0.615094	-0.615094	11/11	11/11
0000+++--+	-0.632973	-0.632973	11/11	11/11
0000--++--	-0.633049	-0.633049	11/11	11/11
0000---+--	-0.709478	-0.709478	11/11	11/11
0000++-++-	-0.709620	-0.709620	11/11	11/11

The repeated value across all 11 double-free LMB subsets is expected for this fixed ray because the ray is supplied in each LMB basis. Orientation 0000+++--+ is the worst fitted one, but it is only slightly worse than 0000--++--; the enhancement is therefore a broad orientation-sector feature rather than a single-orientation defect.

The z-axis exactly collinear equal-radius test ray, $(0,0,1)$ and $(0,0,1)$ with norms 1,1, produced missing fits rather than stable slopes. This was a direction-specific numerical special case, not the generic behavior of coincident rays; the GL00 plot-direction coincident ray gives the stable positive slopes shown above.

These fixed-ray results are not evidence for a generic uncanceled UV divergence. They instead expose a non-uniform, codimension relative-momentum ridge that the random-ray UV profile hides.

This is not in contradiction with the positive-slope event-weight plots. The UV profile uses momentum-space samples with `integrator.weight=1` and multiplies the raw integrand by only the scale-measure factor $s^{3n_{UV}}$. For the double-UV tests here this is s^6 . The radial approach plots, by contrast, show the density seen by the integrator after the spherical linear map. For one hard loop,

$$r = E_{\text{cm}} b \frac{x}{1-x}, \quad r^2 \frac{dr}{dx} \sim r^4,$$

whereas the UV shell-measure factor used by the profiler is r^3 per loop. Thus the plotted x -space weight contains one extra large-radius power per hard loop compared to the profiler's DOD diagnostic. A tailored-ray profiler slope near -0.813 can therefore correspond to a pointwise x -space weight that grows roughly like a positive power along the same radial approach. When the rays are tuned even closer together, the profiler itself finds zero or positive slopes along that forced line. The transverse angular and radius-ratio volume is not represented by this one-dimensional profile, so the positive forced-ray DOD identifies an integrable relative singularity and a variance problem rather than a generic UV divergence of the full integral.

The profiler JSON and logs for the default, plot-ray, simplified split-ray, exactly collinear, and per-orientation tests are saved in the report data directory under the `uv_profile.*` file-names.

7 Current long-run integration checks

The current all-LMB one-million-sample diagnostic, using Monte Carlo over LMB channels, kinematics A, helicity $+-$, $m_{UV} = 9.188$, and $\mu_R = 91.188$, gave:

component	value	error	relative error
real	-6.8192×10^{-6}	4.7659×10^{-6}	69.9%
imag	-5.4228×10^{-6}	3.3372×10^{-6}	61.5%

The largest signed weights in that run were:

component/sign	max evaluation	LMB channel
real +	$+5.3591 \times 10^{-1}$	(5, 9)
real -	-4.3118	(5, 10)
imag +	$+5.4193 \times 10^{-1}$	(5, 9)
imag -	-2.5294	(5, 10)

A longer Monte-Carlo-over-LMB run with 90 M samples and 20 cores completed in 739.3 s. It did not reach the requested sub-10% Monte-Carlo error:

component	value	error	relative error	max-weight impact
real	-4.4150×10^{-7}	4.2964×10^{-7}	97.3%	0.791
imag	-1.2004×10^{-6}	3.0318×10^{-7}	25.3%	0.134

The largest signed weights in this run moved to LMB channel (5, 7). The largest real negative event was -31.4201, and the largest imaginary negative event was -14.4356, both at the same sampled point. The adaptive grid probability for channel (5, 7) was about 42.2%, but the continuous relative ridge still dominated the variance.

An explicit summed-LMB run with 10 M samples was run for a comparable wall time, 740.9 s. Summing removes the discrete LMB sampling noise but makes each sample roughly an order of magnitude more expensive. At equal wall time it also did not converge:

component	value	error	relative error	max-weight impact
real	$+6.6162 \times 10^{-8}$	3.2102×10^{-7}	485.2%	1.513
imag	-1.1342×10^{-6}	4.4162×10^{-7}	38.9%	0.190

The summed-channel run lowers the absolute largest sampled weights compared to the MC-over-LMB run, but the real part is close enough to zero that its relative error remains very large. These results show that the default $\alpha = 3$ LMB multichannel setup did not reach stable sub-10% errors in the tested wall time. They should not be read as proof that LMB multichanneling is irrelevant, or that no choice of α can help; the unresolved issue is the continuous relative-momentum ridge left after the discrete OSE reweighting.

The saved run artifacts are in the report data directory, with separate JSON and log files for the 90 M MC-over-LMB run and the 10 M summed-LMB run.

8 What the LMB multichannel prefactor does

The generated all-LMB setup contains channels where edge 4 is a basis edge, and the LMB multichanneling code does use this information. For a channel c with basis edges $e \in c$, the channel weight is

$$w_c = \left(\prod_{e \in c} E_e \right)^{-\alpha}, \quad p_c = \frac{w_c}{\sum_j w_j}. \quad (12)$$

With `lmb_channels = "monte_carlo"`, one LMB is sampled with this normalized weight. With `lmb_channels = "summed"`, all generated LMBs are evaluated with the same normalized prefactors. The default value is $\alpha = 3$, and the flat `[sampling]` parser now exposes this as an `alpha` field so that it can be scanned directly.

I evaluated this normalized prefactor on the same ray family used in the fixed-ray UV-profile scan. The table reports the raw probability assigned to the problematic (5, 10) interpretation and the total probability assigned to channels containing edge 4:

a	R	α	E_4/R	$p_{(5,10)}$	$\sum_{4 \in c} p_c$
1.000	10^3	1	2.041×10^{-1}	3.050×10^{-2}	0.854066
1.000	10^3	3	2.041×10^{-1}	1.621×10^{-3}	0.992885
1.000	10^7	1	3.566×10^{-2}	5.814×10^{-3}	0.970933
1.000	10^7	3	3.566×10^{-2}	7.726×10^{-6}	0.999961
0.100	10^7	1	3.584×10^{-3}	5.959×10^{-4}	0.997020
0.100	10^7	3	3.584×10^{-3}	7.687×10^{-9}	1.000000
0.020	10^7	1	7.206×10^{-4}	1.200×10^{-4}	0.999400
0.020	10^7	3	7.206×10^{-4}	6.240×10^{-11}	1.000000
0.005	10^7	1	1.843×10^{-4}	3.072×10^{-5}	0.999846
0.005	10^7	3	1.843×10^{-4}	1.044×10^{-12}	1.000000
0.000	10^7	1	1.973×10^{-5}	3.287×10^{-6}	0.999984
0.000	10^7	3	1.973×10^{-5}	1.279×10^{-15}	1.000000

The conclusion is that the OSE prefactor is doing something very strong. On the near-collinear rays, default $\alpha = 3$ assigns essentially all normalized channel weight to LMBs containing edge 4 and suppresses the direct (5, 10) interpretation by many orders of magnitude. Increasing α would make this channel selection even sharper; decreasing it would make the channel selection smoother. Therefore the previous statement that LMB MC merely samples an LMB, or that it cannot be sufficient by construction, would be too strong.

However, the OSE prefactor is not the same object as the balance-heuristic density of the actual channel maps. For channel maps $y = \phi_c(x_c)$ from unit-cube variables to physical loop momenta, a standard multichannel denominator would use

$$p_c(y) \propto \frac{\pi_c}{J_c(y)}, \quad J_c(y) = \left| \det \frac{\partial \phi_c}{\partial x_c} \right|_{x_c = \phi_c^{-1}(y)}, \quad (13)$$

up to the discrete prior π_c . The current code instead computes the partition in Eq. (??): it reinterprets the sampled momenta in the selected LMB, keeps the sampled conformal-map Jacobian, and multiplies by the normalized OSE product. Thus the current denominator knows which LMB energies are small, but it does not know the actual conformal-map density induced by each competing LMB parameterization.

I added a two-channel test to the standalone proxy to isolate this distinction. The two channels are the bad direct interpretation (u, v) and an explicit relative interpretation

$$P_{\text{common}} \simeq \frac{u+v}{2}, \quad q = v - u + Q.$$

For a spherical linear conformal map the one-vector Jacobian used in `parameterize3d` is

$$J_3(r) = 4\pi r^2 \frac{(E_{\text{cm}}b + r)^2}{E_{\text{cm}}b}.$$

The inverse-Jacobian balance denominator was compared to OSE partitions with $\alpha = 1, 3, 4$.

path	raw slope	OSE $\alpha = 1$	OSE $\alpha = 3$	OSE $\alpha = 4$	inverse J , linear
fixed $ q = 0.2$, varying R	+1.00	+0.00	-2.00	-3.00	-3.00
fixed $ q /R = 10^{-3}$	-1.00	-1.00	-1.00	-1.00	-1.23
fixed $R = 10^6$, varying $ q $	-1.16	-0.16	+1.84	+2.84	+0.85

Here “slope” means the fitted slope of the bad direct-channel contribution after multiplying by the corresponding partition factor. On the strict locked-ray path with fixed small q , the raw bad-channel weight grows like R . The default OSE factor with $\alpha = 3$ already changes that contribution to R^{-2} . An inverse-Jacobian denominator for the linear map gives R^{-3} ,

numerically close to OSE $\alpha = 4$. At $R = 10^7$ and $|q| = 0.2$, the bad direct-channel probabilities are

$$p_{\text{bad}}^{\text{OSE}, \alpha=3} = 8.0 \times 10^{-24}, \quad p_{\text{bad}}^{\text{OSE}, \alpha=4} = 1.6 \times 10^{-31}, \quad p_{\text{bad}}^{1/J, \text{linear}} = 3.3 \times 10^{-26}.$$

For the logarithmic radial map, the inverse-Jacobian probability is even smaller in the deep UV because $J_3(r)$ grows exponentially with r/E_{cm} in that map; it underflows below double precision already in the large- R part of this scan. The OSE denominator has no way to know about that map-dependent density.




Figure 9: Standalone proxy comparison of the current OSE partition and an inverse-Jacobian balance partition. The OSE $\alpha = 3$ factor already suppresses the strict small-relative-momentum direct channel strongly. The inverse-Jacobian denominator is more map-faithful and gives one additional power of R for the linear map on the fixed- q locked ray, while the logarithmic map would suppress the bad direct channel even more strongly in the far UV.

8.1 Pointwise inspect check of LMB multichanneling

I also checked the problematic GL00 region directly with `inspect`, using kinematics A and helicity $+-+-$. The current `examples/cli/aa_aa/2L` state is not an all-LMB state: `display integrand` shows only one active LMB channel, channel 0 with basis (5, 10). Therefore, in that state, switching `lmb_multichanneling=true` and `lmb_channels="monte_carlo"` does not actually average over all LMBs; it evaluates the same single selected LMB channel.

For the all-LMB regenerated state, `display integrand` shows eleven active channels:

$$(5, 10), (4, 10), (5, 9), (4, 9), (5, 8), (4, 8), (5, 7), (4, 7), (5, 6), (4, 6), (4, 5).$$

The table below compares the displayed inspect weight, i.e. the integrand times the continuous parameterization Jacobian. These are the weights seen before the additional discrete-channel normalization present in the Monte-Carlo-over-LMB estimator.

setup	point	Re weight	Im weight	weight
LMB off	point A	$-2.63307595 \times 10^{-1}$	$+2.16943924 \times 10^0$	2.18535982×10^0
single active LMB channel 0, (5, 10)	point A	-4.20480647×10^1	$+8.23757741 \times 10^1$	9.24867985×10^1
all LMBs, explicit summed channels	point A	$-1.84215087 \times 10^{-3}$	$+3.52144262 \times 10^{-3}$	$3.97417639 \times 10^{-3}$
LMB off	channel-6 max-weight point	$+7.74559448 \times 10^{-4}$	$-8.69620692 \times 10^{-4}$	$1.16455248 \times 10^{-3}$
LMB channel 6, (5, 7)	channel-6 max-weight point	-3.05497125×10^0	-1.40357390×10^0	3.36197398×10^0

This does not mean that the multichannel denominator only acts when an edge-4 channel is selected. The intended balance mechanism is global at the sampled physical point: even if the sampled chart is (5, 10) or (5, 7), the denominator

$$\sum_c \prod_{e \in c} E_e^{-\alpha}$$

contains the competing edge-4 charts. The question is therefore why this global denominator is not sufficient.

For point A, interpreted in the bad (5, 10) chart, the local normalized prefactor for the selected channel is already small,

$$p_{(5,10)}^{\text{OSE}, \alpha=3} = 8.46 \times 10^{-6}.$$

The relevant edge scales are

$$|e_4| = 1.78 \times 10^3, \quad E_4 = 1.79 \times 10^3, \quad |e_5| = 4.82 \times 10^4, \quad E_{10} = 4.93 \times 10^4.$$

So the OSE denominator is doing the expected qualitative thing in the deep UV approach: the edge-4 charts dominate the denominator and suppress the bad chart. The remaining problem is quantitative and structural: the raw bad-chart density is large enough that this OSE-product suppression does not flatten the estimator, and the OSE product is not the true density of the competing maps.

I then ran a controlled comparison in the same all-LMB regenerated state, with the same seed and 10 M samples:

setup	Re	Im	largest $ w $	top channel
LMB off	$-5.55(8.03) \times 10^{-8}$	$-7.19(0.85) \times 10^{-7}$	2.69×10^{-1}	none
all LMBs, MC over LMBs	$-4.94(7.22) \times 10^{-7}$	$-3.31(10.49) \times 10^{-7}$	7.22	(5, 7)

In this controlled run, and also in the 90 M all-LMB run, the largest weights are in channel (5, 7). This channel is precisely a chart in which the edge-4 momentum is a difference of two sampled basis momenta:

$$e_5 = -K_0 + K_1 + P_0, \quad e_7 = K_1 + P_0, \quad e_4 = K_0 = e_7 - e_5.$$

The necessary edge-4 LMBs are present, for example (4, 7) and (4, 5). What fails is not channel enumeration; it is the feature used in the balance denominator.

At the largest (5, 7) points the edge-4 *three-momentum* is small, but edge 4 is a massive top edge, so its OSE does not become small:

point	$ e_4 $	E_4	selected-channel prefactor	largest sampled channel
10 M all-LMB MC max	8.71	173.22	8.78×10^{-2}	(5, 7)
90 M all-LMB MC max	7.47	173.16	1.17×10^{-1}	(5, 7)

Thus the OSE partition barely sees the small relative three-momentum in this moderate-radius ridge: $E_4 \simeq m_t$ is just a mass scale. This is exactly where an inverse-Jacobian or relative-momentum density would differ from the current OSE proxy. For a spherical linear map, a chart that samples the relative vector $q = e_4$ directly has a Jacobian factor proportional to q^2 near $q = 0$, so its inverse density is enhanced like $1/q^2$ even for a massive edge. The OSE factor $E_4^{-\alpha}$ has no corresponding enhancement because $E_4 = \sqrt{|q|^2 + m_t^2}$.

This answers the inverse-Jacobian question as follows. Yes, replacing the OSE-product denominator by the actual inverse-Jacobian channel density would be more principled and should help in exactly the regimes where the conformal map itself controls the rare-event density, especially for logarithmic mapping and large common radius. It would also make the channel balance automatically adapt to E_{cm} and b . It is also the right diagnostic for the massive-edge issue above, because it is sensitive to the small spatial relative variable even when the corresponding OSE is pinned by the mass. The OSE factor is therefore not useless; it suppresses the deep-UV bad chart by many orders of magnitude. But it is not the same as the balance density of the actual channel maps, and for the massive relative edge it misses the most important local variable, $|e_7 - e_5|$.

This explains why the longer all-LMB Monte-Carlo run assigned large probability to channel (5, 7), about 42.2%, and sizable probabilities to channels such as (5, 10), about 15.8%, while individual edge-4 channels were only about 1.4% each. The adaptive procedure sees the contribution and variance of the existing channels, but it does not infer a new relative-momentum map. The practical recommendation is therefore: disable the LMB-selection heuristic when studying these graphs so that all LMB channels are actually available; keep explicit LMB summation as a diagnostic baseline; and add a continuous relative-momentum channel for hard-hard configurations where an internal momentum such as $k - l + Q$ becomes small while both loop momenta are UV. The multichannel denominator should then use the actual densities of the competing maps, or at least an inverse-Jacobian approximation to them, instead of relying only on the OSE product. An α scan remains useful now that the setting is exposed in `[sampling]`, but tuning α alone cannot create the missing relative variable.

9 Other $\gamma\gamma \rightarrow \gamma\gamma$ two-loop graphs

I reran the full selected two-loop graph set GL00, GL01, GL03, GL05, GL06, GL08, GL12, GL14, GL16, GL17, and GL18. Each graph was generated in a fresh state with integrated UV enabled and a 100 GiB virtual-memory cap. All eleven labels now complete generation and the standard random-ray UV profile. The profiling and integration runs used kinematics A, helicity $+-+-$, disabled thresholds, $\mu_R = 91.188$, and $m_{UV} = 9.188$.

The integration campaign used 20 cores, one graph at a time, summed orientations, generated LMB multichanneling, and Monte-Carlo sampling over LMB channels rather than explicit summation over LMB channels. Each graph was run for a 10-minute wall-time window with `integrator.n.start=100000`, `n.increase=0`, and `n.max=1000000000`. The timeout exit code is therefore expected; the table below uses the final completed iteration written before termination.

graph	feyngen [s]	generation [s]	peak RAM [GiB]	UV entries	worst UV slope	slopes > -0.9
GL00	8	64	13.01	33	-0.99727	0
GL01	16	86	16.32	33	-0.99544	0
GL03	8	74	14.46	33	-0.99575	0
GL05	24	86	13.80	33	-0.99837	0
GL06	7	24	2.83	42	-0.99572	0
GL08	30	42	2.19	45	-0.99796	0
GL12	11	22	1.87	45	-0.99545	0
GL14	2	17	2.77	42	-0.99908	0
GL16	27	39	2.27	45	-0.99940	0
GL17	10	26	2.83	42	-0.99848	0
GL18	6	21	2.32	42	-0.99765	0

All random-ray UV profiles are therefore clean in this diagnostic sense: every fitted slope is negative with comfortable margin, no fitted slope is above -0.9 , and no arbitrary-precision retry was needed. This remains consistent with the GL00-specific observation above: the UV profiler probes generic rays well, while the integrator can still see rare high-weight relative-momentum ridges.

graph	evals [10 ⁶]	avg time [total/int/eval μ s]	Re [10 ⁻⁶]	Im [10 ⁻⁶]	10-minute status
GL00	43.5	158.5/73.5/58.7	-0.991 ± 1.048 (105.8%)	-0.745 ± 1.656 (222.2%)	not converged
GL01	44.4	159.9/73.7/57.0	-2.656 ± 2.092 (78.7%)	1.159 ± 3.721 (321.0%)	not converged
GL03	47.2	145.6/68.1/57.1	0.249 ± 0.703 (282.2%)	1.383 ± 0.932 (67.3%)	not converged
GL05	45.5	150.8/70.3/57.5	0.579 ± 1.673 (289.1%)	2.193 ± 2.122 (96.8%)	not converged
GL06	47.6	129.0/59.5/46.8	-0.008 ± 0.035 (456.6%)	-1.135 ± 0.038 (3.3%)	imaginary part only
GL08	53.8	126.5/54.1/26.3	-0.857 ± 0.846 (98.8%)	-3.435 ± 2.302 (67.0%)	not converged
GL12	58.0	114.0/49.1/26.1	1.476 ± 1.580 (107.0%)	0.505 ± 0.143 (28.2%)	not converged
GL14	43.9	143.7/65.3/47.4	0.009 ± 0.038 (404.3%)	4.234 ± 0.043 (1.0%)	imaginary part only
GL16	49.3	145.8/61.5/26.6	0.325 ± 0.252 (77.4%)	-3.965 ± 0.234 (5.9%)	imaginary part only
GL17	45.3	143.9/65.4/47.3	-0.046 ± 0.040 (87.7%)	4.155 ± 0.044 (1.1%)	imaginary part only
GL18	45.6	141.6/64.4/46.9	-0.033 ± 0.047 (144.6%)	-1.158 ± 0.053 (4.6%)	imaginary part only

No complete graph reached sub-10% Monte-Carlo uncertainty in both real and imaginary parts in the 10-minute, 20-core, LMB-channel-MC setup. The behavior separates into two groups. GL00, GL01, GL03, GL05, GL08, and GL12 remain high-variance in both components. GL06, GL14, GL16, GL17, and GL18 have a stable imaginary component at the percent level, but a real part compatible with zero and therefore a large relative error. Thus the graph-level convergence problem is not only a GL00 issue, although the catastrophic double-UV ridge documented in the earlier sections is the clearest mechanistic example.

The compact campaign summary and per-graph final UV-profile and integration JSON files are archived in the report data directory under `aa_aa_2l_all_graphs_10min_lmbmc`.

9.1 Max-weight geometry and fixed-ray UV profiles

For each graph, I then took the largest absolute entry in the stored `max_weight_info`. The spherical unit-cube coordinates were decoded to the two sampled LMB ray directions and radii. I ran two additional UV profiles from those data:

1. an exactly locked profile, where the two UV rays are made exactly parallel, or anti-parallel when the recorded LMB directions have opposite sign, and are assigned the same starting norm;
2. an actual-angular profile, where the two recorded angular directions are used as-is, again with equal starting norms.

The first profile tests the strict equal-norm relative-momentum ray. The second profile checks whether the finite max-weight point itself already sits on the asymptotic ray. In all cases the usual random-ray UV profile is clean, so the table below reports only the largest fitted slope from each follow-up.

graph	max eval	angle [deg]	r_1/r_0	exact locked slope	actual-angular slope	exact nonnegative fits	interpretation
GL00	-42.9	0.954	1.156	1.000	-1.000	11	confirmed GL00-type ray
GL01	161.5	1.412	0.947	-1.000	-1.000	0	finite near-collinear peak
GL03	-16.3	9.270	1.082	1.000	-1.000	11	confirmed GL00-type ray
GL05	85.0	10.901	0.967	-1.000	-1.000	0	finite near-collinear peak
GL06	-0.536	2.477	0.960	1.000	-0.999	14	confirmed, small residue
GL08	-123.3	0.771	1.016	-0.964	-0.965	0	borderline finite peak
GL12	91.2	179.458	1.040	2.000	-0.962	15	confirmed anti-parallel ray
GL14	0.582	0.434	1.004	1.000	-1.000	14	confirmed, small residue
GL16	7.13	178.023	1.044	2.000	-0.974	15	confirmed anti-parallel ray
GL17	0.812	7.256	0.992	1.000	-1.000	14	confirmed, small residue
GL18	-1.47	179.676	0.985	1.000	-1.000	14	confirmed anti-parallel ray

This confirms that the GL00 diagnostic is not isolated. The same strict equal-norm double-UV ray is exposed by the UV profiler for GL00, GL03, GL06, GL12, GL14, GL16, GL17, and GL18. For GL12, GL16, and GL18, the same condition appears as anti-parallel rays in the sampled LMB basis; this is the same relative-momentum mechanism with a different sign convention in the chosen basis. The positive slopes always occur in the full double-UV subset where both LMB edges are free, for example the `subset=2` entries `free=(5,10)`, `free=(4,10)`, and analogous pairs in the generated LMB list.

The actual-angular fixed profiles are negative even when the recorded max-weight point is less than one degree from collinearity. This is important: the profiler only sees the positive slope when the two UV rays are locked exactly onto the relative-momentum ray. A small opening angle or radius mismatch restores a negative asymptotic slope, while the finite coefficient can still be large enough to dominate a Monte-Carlo iteration.

GL01, GL05, and GL08 are therefore different from the confirmed group in this diagnostic. Their largest weights are still near-collinear and near equal-radius, but the fixed-ray UV profile does not show a nonnegative asymptotic slope. I pushed GL01 and GL05 to `max-scaling=15`; both stayed at slope -1 . GL08 is the borderline case: it gives a mildly degraded double-UV slope near -0.964 at `max-scaling=9` and -0.937 at `max-scaling=15`, but a deeper f128 run to `max-scaling=21` returns to -1.000 . I therefore classify these three as large finite near-collinear peaks rather than confirmed GL00-type positive-DOD rays.

The detailed fixed-ray summary is archived as `aa_aa_2l_all_graphs_10min_lmbmc/max_weight_uv_diagnostic_summary.json`.

10 Implications

The current evidence points to a high-variance, locally integrable relative-momentum singularity rather than a wrong integrated UV counterterm or a missing local UV cancellation. The practical options suggested by the toy model are:

1. add a specialized channel or remapping for hard-hard relative momentum ridges, using variables close to (P_{common}, q) ;
2. scan the LMB prefactor exponent α ; larger values focus more strongly on channels containing the small OSE, while smaller values avoid over-focusing and may improve other regions;
3. use explicit summed LMB channels as a diagnostic, while recognizing that this reduces discrete-channel variance but does not by itself remap the continuous relative-ridge variables.